

TTC'15 Case Study: Transformation of Java Annotations

Filip Křikava

INRIA Lille

France

`filip.krikava@inria.fr`

Martin Monperrus

University Lille 1 / INRIA Lille

France

`martin.monperrus@univ-lille1.fr`

Java 5 introduced annotations as a systematic mean to attach syntactic meta-data to various elements of Java source code. Since then, annotations have been extensively used by a number of libraries, frameworks and tools to conveniently extend behavior of Java programs that would otherwise have to be done manually or synthesized from external resources. The annotations are usually processed through reflection and the extended behavior is injected into Java classes using aspect-oriented techniques or a direct byte code modification. However, in some cases, class-level instrumentation might not always be available neither desirable and therefore the transformation is done at the source code level.

In this case study we focus on such source-level transformation. Concretely, the task is to inject behavior specified by an annotation library that encapsulates common programming concerns such as logging, caching and a failure retry. The objective is to explore how are the contemporary transformation tools suitable for programming language transformations.

1 Introduction

Java 5 has been extended with annotations that provide a convenient way to supply meta-data to various element of Java source code such as classes and methods. Since then, annotations have been extensively used by a number of Java libraries, frameworks and tools. One of the advantage of using annotations is that they attach syntactic meta-data directly to the relevant element of Java source code rather than decoupling them to external resources or setting them up manually through some API.

Among other things, annotations can be used to express crosscutting concerns such as logging or caching that would otherwise cluttered the source code as they are difficult to encapsulate using regular programming techniques [3]. Essentially, an annotation acts as a marker for a transformation tool which applies the expressed concern into the corresponding Java element (*e.g.*, a class, a method). Majority of these tools rely on Java reflection API and some sort of class instrumentation using aspect-oriented programming (AOP) techniques or direct byte code manipulation to process the these annotations. While this is usually the preferred way, in some cases, the class level instrumentation might not available neither desirable. The problem is that it might (i) lead to leaky abstraction since the transformation is transparent to a developer, (ii) introduce some performance penalties, or (iii) bring some extra non-trivial dependencies.

In this case study we focus on the situation when class level instrumentation is not possible and develop a solution for annotation processing based on source level transformation. Concretely,

the transformation task is to inject behavior specified by an annotation library that covers common programming concerns such as logging, caching and retry certain type of failures. The objective is to explore how the current model transformation tools are suitable for programming language transformations; particularly in comparison to some dedicated program transformation libraries such as Spoon [4]. The idea to bring together different communities on a common transformation problem.

In the next section we present the sample annotation library and describe the tasks of the case study in more details. After that, Section 3 overviews criteria that will be used to evaluate submitted solutions.

All supporting material, as well as this document, is available on the case study [github page](#)¹. The contestants are invited to submit an issue shall there be any ambiguity about the transformation tasks that are carried in this case study.

2 Case Description

As the annotation library, we use a simplified version of jcabi-aspects [2], a collection of Java annotations together with AOP-based processing tool that allows developers to conveniently apply some common crosscutting concerns into Java applications. In order to balance the development effort in this case study, we have selected the following annotations from the jcabi-aspects library:

- `@RetryOnFailure` for an automated retry of failed method execution,
- `@Cacheable` for a simple data caching of parameterless methods return values, and
- `@Loggable` for an automated logging of method calls.

The core transformation involves traversing Java source code and extending the annotated methods with the behavior specified by these annotations. The details of each annotation including the specification of its arguments are provided in its associated [javadoc](#).

In the following text we provide an overview of the individual transformation tasks that are part of this case study. The transformations are illustrated on an example rather than formally described which we believe should be sufficient² given the nature of the transformation. The first three tasks consider each of the listed annotation individually, while in the last task, we look on the case of combining these annotations together. Finally, we also propose an extension for an optimization the resulting Java source code when these annotations are combined. The complete code for these transformation tasks is provide on the github page in the [ttc15.tranj.examples package](#).

2.1 Running Example

To better illustrate the transformation tasks of this case study, we use the following example. Let us consider a class that is used to download a content of an URL. A possible³ Java implementation

¹<https://github.com/fikovnik/ttc15-tranj-case>

²Shall there be any ambiguity, the readers are invited to raise an issue on the case study [github page](#).

³In order to keep the case study concise and clear we omit the use of any external library.

is shown in Listing 1.

```

1  public class URLDownload {
2      private final URL url;
3
4      public URLDownload(String url) throws MalformedURLException {
5          this.url = new URL(url);
6      }
7
8      public byte[] get() throws IOException {
9          try (InputStream input = url.openStream()) {
10
11              ByteArrayOutputStream buffer = new ByteArrayOutputStream();
12              byte[] chunk = new byte[4*1024];
13              int n;
14
15              while ((n = input.read(chunk)) > 0 ) {
16                  buffer.write(chunk, 0, n);
17              }
18
19              return buffer.toByteArray();
20          }
21      }
22  }

```

Listing 1: Basic version of `URLDownload` class.

When the basic functionality is implemented, we would like to extend it with the following features:

- *Failure handling.* The method should become more robust and accommodate for some of the inevitable network delays by retrying the download in the case an exception occurs. It should, however, only retry the call in the case when it makes sense—*i.e.* when it is possible to recover. For example, in the case of `SocketTimeoutException` exception, the execution should be retried (preferably after a delay), while `UnknownHostException` it should fail immediately.
- *Caching.* The URL should not be consulted every single time the method is called, but instead it should be keep the result of the last invocation in memory and reuse it for certain amount of time.
- *Logging.* Each call to the method should be logged. The logging message should present the state of the current instance at the method entrance as well as how long the invocation took at the method exit. All exceptions handled within the method body should be also logged.

Instead of coding this manually, producing code such as the one shown in the listing in Section A, we would like to use annotations to declaratively specify the above concerns, and have a transformation tool automatically synthesize code similar to the listing shown in Section B. Concretely, the only change to the original code should be the following annotations:

```

1  @RetryOnFailure(attempts = 3, delay = 1000, retry = { SocketTimeoutException.class },
2      ↪ escalate = { UnknownHostException.class })
3  @Cacheable(lifetime = 1000)
4  @Loggable
5  public byte[] get() throws IOException { ... }

```

2.2 Task 1: Retrying on a Failure

The first transformation tasks should handle the `@RetryOnFailure` annotation. The objective is to extend a given method with a simple failure handling strategy that retries failed invocations according to given options. Annotating the `URLDownload.get()` method with

```
1 @RetryOnFailure(attempts = 3, delay = 1000, retry = { SocketTimeoutException.class },
   ↪  escalate = { UnknownHostException.class })
```

should produce the following code:

```
1 public byte[] get() throws IOException {
2     // added retry counter
3     int __retryCount = 0;
4
5     // added loop
6     while (true) {
7         try {
8             // the content of the original method
9             ...
10        } catch (UnknownHostException e) {
11            // added escalation cases
12            throw e;
13        } catch (SocketTimeoutException e) {
14            // added retry cases
15            __retryCount += 1;
16
17            if (__retryCount > 3) {
18                throw e;
19            } else {
20                // added delay
21                try {
22                    Thread.sleep(1000);
23                } catch (InterruptedException e1) {
24                    throw e;
25                }
26            }
27        }
28    }
29 }
```

The original method code—*i.e.* what is between the lines 9–20 in Listing 1, is wrapped by another try-catch block (line 7) which is itself in a while loop (line 6). The number of attempts (3) is projected in the condition `if (__retryCount > 3)` (line 17) and the delay is translated in a current thread sleep on lines 21–25. Finally, we make a distinction in the exceptions that on which the call should be retried (the `retry` parameter – `SocketTimeoutException`) and the ones that should be escalated (the `escalate` parameter – `UnknownHostException`) in the generated catch clauses in line 10 and 13 respectively.

2.3 Task 2: Caching

The second transformation task concerns the `@Cacheable` annotation. The objective is to extend a given method with a simple caching strategy that keeps a result of a method invocation for a

given period of time. Annotating the `URLConnection.get()` method with `@Cacheable(lifetime = 1000)` should produce the following code⁴:

```

1 public class URLDownload {
2     // added bookkeeping fields
3     private long __getCacheLastAccessed = 0;
4     private byte[] __getCacheContent = null;
5
6     ...
7
8     public byte[] get() throws IOException {
9         // added condition
10        if (System.currentTimeMillis() - __getCacheLastAccessed < 1000 && __getCacheContent
11            ↳ != null) {
12            return __getCacheContent;
13        }
14        try (InputStream input = url.openStream()) {
15            ByteArrayOutputStream buffer = new ByteArrayOutputStream();
16            byte[] chunk = new byte[4 * 1024];
17            int n;
18
19            while ((n = input.read(chunk)) > 0) {
20                buffer.write(chunk, 0, n);
21            }
22
23            // added
24            __getCacheContent = buffer.toByteArray();
25            __getCacheLastAccessed = System.currentTimeMillis();
26
27            // added
28            return __getCacheContent;
29        }
30    }
31 }

```

The `__getCacheLastAccessed` and `__getCacheContent` defined on line 3 and 4 are cache bookkeeping fields used for storing the information about the last access time and method result respectively. The condition on line 10 checks the validity of the cache based on the annotation `lifetime`. Finally, the method exit point is replaced with the stored result on lines 24–25. Clearly, this have to be done for all method exit points in the case the annotated method have multiple ones.

2.4 Task 3: Logging

The last transformation task introduces logging. Any method annotated with `@Loggable` should have conditionally (depending on the annotation options) logged (i) its entry point, (ii) all its exit points, and (iii) all exceptions that are caught within the method body.

Annotating the `URLConnection.get()` method with `@Loggable` should produce the following code:

```

1 public class URLDownload {
2     // added logger
3     private final org.slf4j.Logger __logger =
4         ↳ org.slf4j.LoggerFactory.getLogger(URLDownload.class);
5
6     ...

```

⁴Too keep things simple, we do not concern ourselves with invalidating the cache and thus freeing occupied memory.

```

7  public byte[] get() throws IOException {
8      // added entry point logging
9      long __entryTime = System.currentTimeMillis();
10     if (__logger.isTraceEnabled()) {
11         __logger.trace(String.format("get() [url='%s']: entry", url));
12     }
13
14     try (InputStream input = url.openStream()) {
15
16         ByteArrayOutputStream buffer = new ByteArrayOutputStream();
17         byte[] chunk = new byte[4 * 1024];
18         int n;
19
20         while ((n = input.read(chunk)) > 0) {
21             buffer.write(chunk, 0, n);
22         }
23
24         // added exit point logging
25         if (__logger.isTraceEnabled()) {
26             __logger.trace(String.format("get(): exit in %d ms", System.currentTimeMillis() -
27                                     ↪ __entryTime));
28         }
29         return buffer.toByteArray();
30     }
31 }

```

Line 3 defines an instance of SLF4J logger [1]. Lines 9–12 and 25–27 insert the entry and exit point logging respectively. The entry point logging message additionally includes the state of the object—*i.e.* the value of all its declared (non-inherited) fields.

Because, there is no exception handling code—*i.e.* no catch clause, only entry and exit points are traced. If we would, however, apply the logging annotation on the method which is the output of `@RetryOnFailure` transformation, log messages should be also added right after every catch clause. The result should therefore look like:

```

1  while (true) {
2      try {
3          ...
4      } catch (UnknownHostException e) {
5          // added exception handling
6          __logger.error("get(): exception", e);
7
8          throw e;
9      } catch (SocketTimeoutException e) {
10         // added exception handling
11         __logger.error("get(): exception", e);
12
13         __retryCount += 1;
14
15         if (__retryCount > 3) {
16             throw e;
17         } else {
18             try {
19                 Thread.sleep(1000);
20             } catch (InterruptedException e1) {
21                 // added exception handling
22                 __logger.error("get(): exception", e);
23
24                 throw e;
25             }
26         }
27     }
28 }

```

2.5 Task 4: Annotation composition

It should be possible to have all the extension present together on one method:

```
1 @RetryOnFailure(attempts = 3, delay = 1000, escalate = { UnknownHostException.class })
2 @Cacheable(lifetime = 1000)
3 @Loggable
4 public byte[] get() throws IOException { ... }
```

The code transformation should be performed in the declared order—*i.e.*, first applying `@RetryOnFailure`, then `@Cacheable` and finally adding `@Loggable`. The expected result is shown in Section B.

2.6 Extension 1: Optimization (*optional*)

An optional extension to this case study is to have the annotations being aware of one another. In this case, the order of the annotations does not matter and they will be always applied in the order specified in the task 4. For the caching and retry on failure, nothing will change, but the logging will recognize the generated code and it will adjust the logging messages appropriately. The expected result—*i.e.* placement and message content is shown in Section A.

3 Evaluation Criteria

The submitted solutions will be evaluated on a synthetic as well as on a real source code. The transformation is therefore expected to be packed in the way that it is runnable from a command line shell. The program should take two arguments, a path to a *valid* Java 7 source file input and a path to where the result transformation should be stored.

```
$ ./my-solution.sh URLDownload.java SynthesizedURLDownload.java
```

A part of the evaluation process is done automatically and therefore the ability of scripting the transformation is important. Furthermore, it also demonstrates the usability of the solution. The evaluation test cases will be made available to the contestants after the submission period allowing them to reevaluate theirs and others solutions. For this case study we have chosen the following criteria that are evaluated automatically:

- *Correctness.* The transformation tasks are scored individually on the 0 – 1 scale that corresponds to the percentage of successful transformations—*i.e.* the number of correctly inserted code versus the total number of expected insertion points. As a prerequisite, the resulting file must be a *valid* Java 7 file that compiles and whose original behavior is not altered.
- *Performance.* All submissions will be also assessed on their performance—*i.e.* how long did it take to transform the individual test cases.

Next to the automatic evaluation, we are also interested in the assessment of the usability of a given transformation tool to the problem being addressed. Usability of a programming language, a library or a tool is difficult to assess as it tends to be subjective since it largely depends on the preferences and background of its users. In this case study we use a combination of source code complexity, abstraction level and development effort. For each task defined in the previous section, we ask the solution authors to provide a measures of

- The number of source lines of code (SLOC) of all manually written code excluding comments and empty lines.
- The time in person-hours spent in developing the solution including design, implementation, testing and debugging.

Interpreting SLOC metrics is always problematic and the issue of what is the right level of “*verbosity*” in a language is complex and should not be reduced naively to just counting SLOC. On the one hand, usability is not achieved only by having fewer lines of code, but instead, by having more expressive and concise code, which is beneficial to writers as well as to readers. On the other hand, code bloat resulting from code duplication and from lack of constructs that enable the building of more concise but expressive statements, is not desirable.

Next to this, we also ask for an assessment about the *abstraction level* that is provided by the transformation tool to solve this case study. The level of abstraction is a degree of the expressiveness of the solution ranging from high to low. The objective is to determine how closely the solution expresses the problem being addressed. A high abstraction level has the solution expresses in the terms of problem-level concepts while a low-level abstraction level is using simple implementation-level concepts.

References

- [1] ASF: *Simple Logging Facade for Java*. Available at <http://www.slf4j.org/>.
- [2] Yegor Bugayenko: *Useful AOP aspects, incl. JSR-303*. Available at <http://aspects.jcabi.com/index.html>.
- [3] Gregor Kiczales, John Lamping, Anurag Mendhekar, Chris Maeda, Cristina Lopes, Jean-Marc Loingtier & John Irwin (1997): *Aspect-oriented programming*. In: *ECOOP'97—Object-oriented programming*, Springer, pp. 220–242.
- [4] Renaud Pawlak, Martin Monperrus, Nicolas Petitprez, Carlos Noguera & Lionel Seinturier (2014): *Spoon v2: Large Scale Source Code Analysis and Transformation for Java*. Technical Report hal-01078532, Inria. Available at <https://hal.inria.fr/hal-01078532>.

A Source Code of URLDownloader Coded Manually

```

1  package ttl15.tranj.examples;
2
3  import java.io.ByteArrayOutputStream;
4  import java.io.IOException;
5  import java.io.InputStream;
6  import java.net.MalformedURLException;
7  import java.net.SocketTimeoutException;
8  import java.net.URL;
9  import java.net.UnknownHostException;
10
11 import org.slf4j.Logger;
12 import org.slf4j.LoggerFactory;
13
14 public class ManualURLDownload {
15     private static final long CACHE_TIMEOUT = 1000;
16     private static final int MAX_RETRY = 3;
17
18     private final Logger logger = LoggerFactory.getLogger(ManualURLDownload.class);
19
20     private long lastAccessed = 0;
21     private byte[] cachedContent = null;
22     private final URL url;
23
24     public ManualURLDownload(String url) throws MalformedURLException {
25         this.url = new URL(url);
26     }
27
28     public byte[] get() throws IOException {
29         long entryTime = System.currentTimeMillis();
30
31         if (logger.isTraceEnabled()) {
32             logger.trace(String.format("get() [url='%s']: entry", url));
33         }
34
35         if (System.currentTimeMillis() - lastAccessed < CACHE_TIMEOUT && cachedContent !=
36             ↪ null) {
37             if (logger.isTraceEnabled()) {
38                 logger.trace(String.format("get(): exit cached [%d ms]",
39                     ↪ System.currentTimeMillis() - entryTime));
40             }
41             return cachedContent;
42         }
43
44         int retryCount = 0;
45
46         while (true) {
47             try (InputStream input = url.openStream()) {
48                 ByteArrayOutputStream buffer = new ByteArrayOutputStream();
49                 byte[] chunk = new byte[4*1024];
50                 int n;
51
52                 while ((n = input.read(chunk)) > 0) {
53                     buffer.write(chunk, 0, n);
54                 }

```

```

55     cachedContent = buffer.toByteArray();
56     lastAccessed = System.currentTimeMillis();
57
58     if (logger.isTraceEnabled()) {
59         logger.trace(String.format("get(): exit [%d ms]", System.currentTimeMillis() -
60             ↪ entryTime));
61     }
62     return cachedContent;
63 } catch (UnknownHostException e) {
64     logger.error("get(): exception -> forced exit", e);
65     throw e;
66 } catch (SocketTimeoutException e) {
67     if (logger.isWarnEnabled()) {
68         logger.warn("get(): exception -> retrying", e);
69     }
70
71     retryCount += 1;
72
73     if (retryCount > MAX_RETRY) {
74         if (logger.isTraceEnabled()) {
75             logger.trace("get(): max retry count reached");
76         }
77         throw e;
78     } else {
79         try {
80             Thread.sleep(1000);
81         } catch (InterruptedException e1) {
82             if (logger.isWarnEnabled()) {
83                 logger.warn("get(): interrupted while waiting to retry");
84             }
85             throw e;
86         }
87         if (logger.isTraceEnabled()) {
88             logger.trace("get(): retrying");
89         }
90     }
91 }
92 }
93 }

```

B Source Code of Synthesized URLDownloader

```

1  package ttcl5.tranj.examples;
2
3  import java.io.ByteArrayOutputStream;
4  import java.io.IOException;
5  import java.io.InputStream;
6  import java.net.MalformedURLException;
7  import java.net.SocketTimeoutException;
8  import java.net.URL;
9  import java.net.UnknownHostException;
10
11 import org.slf4j.Logger;
12 import org.slf4j.LoggerFactory;
13

```

```

14 public class SynthesizedURLDownload {
15     private final Logger __logger = LoggerFactory.getLogger(SynthesizedURLDownload.class);
16     private long __getCacheLastAccessed = 0;
17     private byte[] __getCachedContent = null;
18
19     private final URL url;
20
21     public SynthesizedURLDownload(String url) throws MalformedURLException {
22         this.url = new URL(url);
23     }
24
25     public byte[] get() throws IOException {
26         long __entryTime = System.currentTimeMillis();
27         if (__logger.isTraceEnabled()) {
28             __logger.trace(String.format("get() [url='%s']: entry", url));
29         }
30
31         if (System.currentTimeMillis() - __getCacheLastAccessed < 1000 && __getCachedContent
32             ↪ != null) {
33             if (__logger.isTraceEnabled()) {
34                 __logger.trace(String.format("get(): exit [%d ms]", System.currentTimeMillis() -
35                     ↪ __entryTime));
36             }
37             return __getCachedContent;
38         }
39
40         int __retryCount = 0;
41         while (true) {
42             try (InputStream input = url.openStream()) {
43
44                 ByteArrayOutputStream buffer = new ByteArrayOutputStream();
45                 byte[] chunk = new byte[4*1024];
46                 int n;
47
48                 while ((n = input.read(chunk)) > 0) {
49                     buffer.write(chunk, 0, n);
50                 }
51
52                 __getCachedContent = buffer.toByteArray();
53                 __getCacheLastAccessed = System.currentTimeMillis();
54
55                 if (__logger.isTraceEnabled()) {
56                     __logger.trace(String.format("get(): exit [%d ms]", System.currentTimeMillis()
57                         ↪ - __entryTime));
58                 }
59                 return __getCachedContent;
60             } catch (UnknownHostException e) {
61                 __logger.error("get(): exception", e);
62                 if (__logger.isTraceEnabled()) {
63                     __logger.trace(String.format("get(): exit [%d ms]", System.currentTimeMillis()
64                         ↪ - __entryTime));
65                 }
66                 throw e;
67             } catch (SocketTimeoutException e) {
68                 __logger.error("get(): exception", e);
69                 __retryCount += 1;
70             }
71         }
72     }
73 }

```

```
68     if (__retryCount > 3) {
69         if (__logger.isTraceEnabled()) {
70             __logger.trace(String.format("get(): exit [%d ms]",
71                                     ↪ System.currentTimeMillis() - __entryTime));
72         }
73         throw e;
74     } else {
75         try {
76             Thread.sleep(1000);
77         } catch (InterruptedException e1) {
78             __logger.error("get(): exception", e);
79             if (__logger.isTraceEnabled()) {
80                 __logger.trace(String.format("get(): exit [%d ms]",
81                                         ↪ System.currentTimeMillis() - __entryTime));
82             }
83             throw e;
84         }
85     }
86 }
87 }
```